
GORGO: Online Tuning for Cross-Region Network-Aware LLM Serving

Anonymous Author(s)

Affiliation

Address

email

Abstract

1 Increasingly, LLM inference services proxy client requests to engine replicas distributed globally. Load-balancing policies must jointly account for factors including KV-cache locality, replica load, and variable network latency when optimizing for metrics like latency and TTFT. However, existing systems only evaluate a subset of these factors in their cost model, leading to uneven concentrations of load and KV-cache across replicas. We present GORGO, a proxy architecture that holistically factors network latency, prefill cost, and queueing delay using tunable parameters. Since open-source chat datasets such as LMSYS-Chat-1M and WildChat-4.8M lack long-context, high prefix-reuse data, we release a synthetic dataset, ARTChat-411k, from long-context production metadata. On a tuning window from ARTChat-411k, evolutionary strategies guide the GORGO policy’s parameters to directly optimize p95 TTFT. During held-out evaluation windows, we fix the parameter values learned from tuning and improve p95 TTFT by 8–15% and end-to-end (E2E) latency by 15–31% over baseline load balancing policies simple session affinity and prefix-cache. Code can be found at <https://github.com/Arcadia-Research-Team/GORGO>.

17 1 Introduction

18 In LLM serving systems, perceived latency to the user is dominated by the time-to-first-token (TTFT). On a single replica, TTFT is dominated by three costs: prefill time, round trip time (RTT) from client (proxy) to replica, and queueing delay behind in-flight requests. Prefix-caching, which is enabled in modern inference engines SGLang [Zheng et al., 2024b] and vLLM [Kwon et al., 2023], eliminates the prefill cost of previous turns in a multi-turn conversation. As LLM context windows increase in length, the time saved by prefix caching 90% of a prompt with 100,000 tokens reduces the prefill cost to 10,000 tokens and decreases TTFT substantially.

25 Since modern LLM deployments proxy requests to inference engines across regions, the cost savings of prefix-caching depend on choosing a replica with the request session’s prefix. Popular routing policies such as consistent hashing and prefix reuse aim to distribute load evenly while creating affinity between a session’s requests and replica(s) to maximize KV-cache reuse [Karger et al., 1997, Stoica et al., 2003, Zheng et al., 2024b, Xia et al., 2025]. In compute-constrained regimes, bursty workloads can saturate a high-affinity replica, causing head-of-line (HOL) queueing delays from decode memory contention and negating cost savings from prefix-cache reuse. Routing policies that holistically evaluate all costs related to TTFT can maintain high prefix-cache reuse while minimizing the negative effects of load saturation and heterogeneous network latency.

34 Existing load balancing policies such as least-load, session affinity, and prefix-reuse may account for replica load or KV-cache hit rate; however, no existing policies consider network latency in cross-region scenarios, which can range on the order of 10ms to 1s (Figure 2). GORGO’s routing policy

accounts for all three TTFT costs, normalizes the units of measurement via tunable parameters, and jointly optimizes parameters through online tuning on real user workloads. To tune and stress-test different routing policies, in (§3) we compile a LLM traffic trace from real production requests with high prefix-reuse and long-context prompts. The trace follows Mooncake’s FAST’25 format [Qin et al., 2024] containing per-request timestamps, which can be linearly scaled to simulate variable saturation profiles.

We benchmark GORG0 on a series of user workloads from ARTChat-411k, our sensitized production Mooncake trace, and sweep across variable time scales to effectively saturate replicas without simulating unrealistic HOL queueing delay. Over existing load balancing policy baselines, GORG0 jointly balances optimal TTFT with request concentration across replicas. Under the continuous batching paradigm, the ES-driven hillclimb tuner exploits warmer replicas close to the proxy and dramatically reduces TTFT at the cost of end-to-end (E2E) latency. We map the pareto frontier of request concentration across replicas, which explodes E2E latency for unbalanced distributions, and TTFT. Our contributions help contextualize the performance of LLM proxy routing policies in real-world user workloads, and (§5) lists a case-by-case scenario of when one would want to use aforementioned baseline policies, the online GORG0 policy, and offline GORG0 with held-out weights. Finally, we show how conditioning the GORG0 cost model on proxy-recorded replica load mitigates subversion of TTFT via continuous batching and slashes request latency across a panoply of metrics.

2 GORG0

2.1 Cost Model

TTFT can be decomposed into three different costs: network latency, prefill time, and queueing delay [He et al., 2025]. The KV-Cache, which stores key-value pairs of input sequences, removes redundant prefill computation for user sequences sharing a prefix with previously computed sequences [Zheng et al., 2024b]. In a cross-region deployment setting, for an inference engine replica i holding a set of cached token prefixes c_i , the cost of TTFT can be defined as the following function, where x_r is the input sequence of tokens for a request r and the prefill time depends only on the set difference $(x_r \setminus c_i)$, the tokens in x_r not already cached on i :

$$\text{TTFT} = T_{\text{network}}(i) + T_{\text{queue}}(i) + T_{\text{prefill}}(x_r \setminus c_i) \quad (1)$$

Theoretically, T_{network} and T_{queue} correspond with round trip time (RTT) from a client to server and the duration of request processing ahead of the incoming request r . However, modern inference engines support batching requests continuously in order to minimize waiting for completion of previous request processing [Yu et al., 2022]. While continuous batching allows admission of a new request into the currently running batch, the batch size is still bounded by a maximum number of concurrent requests and the available KV-cache budget, a critical knob that governs how much load a replica can admit before further requests wait in a queue [Kwon et al., 2023].

In a distributed system, the temporal delay of retrieving queueing metrics from an engine replica makes cost evaluation challenging. We represent the input to T_{queue} as the total number of tokens of requests without a completion event on the client proxy. T_{network} is measured trivially as the exponential weighted moving average of a ping’s round trip time from client proxy to server.

$$\text{TTFT} = T_{\text{network}}(\text{RTT}_i) + T_{\text{queue}}\left(i, \sum_{j=1}^{n_r} x_j\right) + T_{\text{prefill}}(x_r \setminus c_i) \quad (2)$$

2.2 GORG0 Proxy Design

The GORG0 proxy routes client requests to the engine replica with the minimum calculated cost from Equation 3, parameterized by weights W_{rtt} , W_{prefill} , and W_{queue} . These parameters weight the inputs T_{network} , T_{queue} , and T_{prefill} , which unfairly mixes the units of time and tokens, to normalize the correlated costs of latency, prefill cost, and queueing time on TTFT. The weight of

Table 1: Workload characterization. Token counts use the same byte-pair tokenizer in all rows. Reuse percentages are token-weighted prefix-trie savings as described in §3.

Dataset	Total reqs	Users	Avg input tokens	Intra-user reuse	Global reuse
LMSYS-Chat-1M	1,000,000	—	467	—	9.0%
WildChat-4.8M	3,199,860	1,833,730	2,925	5.3%	34.4%
ARTChat-411k	411,169	4,984	21,043	53.7%	55.3%

81 T_{prefill} is fixed to 1 because GORGO proxy makes routing decisions based on relative replica cost:
 82 only the ratio of weights matters in this design.

$$\text{TTFT} = W_{\text{rtt}} * T_{\text{network}}(\text{RTT}_i) + W_{\text{queue}} * T_{\text{queue}} \left(i, \sum_{j=1}^{n_r} x_j \right) + T_{\text{prefill}}(x_r \setminus c_i) \quad (3)$$

83 GORGO proxy uses a simple (1+1) evolutionary strategy to tune weights W_{rtt} and W_{queue} on the
 84 objective function, P95 TTFT. Each parent weight $x_{t,k}$ is perturbed multiplicatively in log-space by
 85 a normal random variable z_k times step size σ , and the new offspring weight x'_k is clamped to values
 86 in the hyperparameter range $[lo_k, hi_k]$ where $lo_k > 0$ and $hi_k > 0$.

$$x'_k = \text{clip}(\exp(\ln(x_{t,k}) + \sigma_t * z_k), [lo_k, hi_k]) \quad (4)$$

87 When x' beats the parent weight x_t on the objective metric, the incumbent weight x_{t+1} is updated
 88 to x' , and σ is adjusted to maintain Rechenberg’s 1/5 success rate of 1 accepted offspring for every
 89 5 offspring [Rechenberg, 1973].

90 3 Workload Characterization

91 Routing policies that perform well on short, low-reuse prompts may perform differently on long,
 92 high-reuse prompts because the prefill term in TTFT scales with uncached input tokens. We measure
 93 three datasets along the same axes: request length, conversational structure, and token-weighted
 94 prefix reuse.

95 **Datasets.** **LMSYS-Chat-1M** [Zheng et al., 2024a] is a dataset of one million chatbot conversa-
 96 tions from a publicly hosted demo of a Vicuna model and the ChatbotArena website. **WildChat-**
 97 **4.8M** [Zhao et al., 2024] is a corpus of 3.2M ChatGPT-proxy conversations grouped by client IP.
 98 Our **ARTChat-411k trace** contains chat completion requests from a production GLM 5.1 model
 99 endpoint hosted by an industry company [GLM Team, 2024]. It contains 411,169 requests across
 100 4,984 distinct users, yielding ~ 82 requests per user on average. We used the request metadata for
 101 real-world timestamps and tokenized request bodies for reuse calculation.

102 We measure prefix reuse with a token-weighted prefix trie that ingests each request’s tokenized mes-
 103 sage sequence in arrival order, partitioned by user where the dataset carries user identity. The trie
 104 reports *intra-user savings* (fraction of tokens matching a prefix from the same user’s prior requests),
 105 *cross-user savings* (additional fraction from pooling across users), and *global savings* (sum). LM-
 106 SYS does not carry user identity, so we report only the global figure.

107 Three contrasts stand out (Table 1, Figure 1). *Request length*: ARTChat-411k prompts are $45\times$
 108 longer than LMSYS and $7\times$ longer than WildChat. At typical rates, prefill on a 21,000-token un-
 109 cached prompt takes on the order of seconds.

110 *User concentration*: ARTChat-411k has 82 requests per user on average; WildChat has 1.7; LMSYS
 111 does not carry user identifiers, so intra-user statistics cannot be computed. *Reuse magnitude*: 53.7%
 112 of ARTChat-411k tokens are part of a prefix seen earlier from the same user (5.3% for WildChat,
 113 essentially zero for LMSYS). The bulk of WildChat’s 34.4% global reuse comes from cross-user
 114 shared structure (chat templates, viral prompts), not sustained per-user dialogue. Thus, for routing,
 115 LMSYS and WildChat live in a regime where prefix-aware policies have little to win. On ARTChat-
 116 411k, a strategy that lands a user on a replica that has seen their conversation can save more than
 117 half of the prefill.




Figure 1: *Left*: Token composition per dataset. ARTChat-411k’s reuse is overwhelmingly intra-user (54%), reflecting multi-turn dialogue; WildChat’s is predominantly cross-user (29%, shared templates); LMSYS has minimal reuse (9%, all cross-user). *Right*: Radar chart with six axes normalized to the dataset maximum. ARTChat-411k dominates on every axis relevant to routing: long prompts, high multi-turn density, concentrated users, and strong intra-user prefix reuse.

Experimental windows. We tune on a high-diversity 30-minute window from ARTChat-411k traffic on April 5th, 2026: 16:15–16:45 UTC (W1). We then freeze the learned parameters and evaluate on two non-consecutive, held-out windows drawn from later days and time-of-day regimes: April 6th, 15:05–15:35 UTC (W2a), and April 7th, 19:45–20:15 UTC (W2b). This split tests whether parameters learned from one high-diversity window generalize to later traffic rather than simply carrying over to an adjacent interval. The two evaluation windows were selected from the 7-day trace for high user diversity and multi-turn density while avoiding single-user-dominated bursts. Traces are replayed at original arrival speed and pre-filtered to 24,000 input tokens. Output is capped at 128 tokens to prevent decode from dominating latency.

4 Experimental Setup

Hardware and software. Each policy runs on its own dedicated three-replica fleet so routing decisions cannot warm or cool another policy’s caches. Each fleet consists of one $2 \times \text{L40S SGLang}$ [Zheng et al., 2024b] server (tensor-parallel 2, 96 GB VRAM per replica) in each of Asia (Seoul), Europe (Frankfurt), and US East (Ashburn), serving Qwen3.5-35B-A3B-FP8 [Qwen Team, 2025]. The model’s native context window is 32,768 tokens; we cap workload input at 24,000 tokens to leave KV headroom for in-flight batches. A CPU proxy in US East (Ashburn) hosts the routing policy and replays the trace.

Cross-region RTT distribution. Probed from the US East proxy over the Apr 5 16:15–16:45 tuning window, EWMA-smoothed RTT to the US East (Ashburn) replica is 30–45 ms (mean 37 ± 8 ms), to Frankfurt 200–390 ms (mean 279 ± 80 ms), and to Seoul 260–700 ms (mean 456 ± 138 ms). The wide Seoul swing reflects routing-table churn along the trans-Pacific path during the trace, while Ashburn stays near-constant. Mean RTT spans a $12 \times$ range across regions, widening to $\sim 23 \times$ at the Seoul peak. Figure 2 shows the full RTT time series over the tuning window.

Baselines, policies, and windows. We compare the three GORGO variants of §2 (gorgo-static, gorgo-autotune, gorgo-hillclimb) against five baseline policies. *[Reviewer comment: Confirmed: we are moving away from gorgo-autotune. It does not appear in the main 2D results (Table 2); it survives only in the WildChat appendix (Table 7), where it was the unstable rate-inversion variant. If dropping it, remove it here and from Appendix A.]* random samples a replica uniformly per request and serves as the lower bound. least-request routes to the replica with the fewest in-flight requests, taking the larger of the proxy’s in-flight view and SGLang’s scraped running-request count to bridge the ~ 10 s staleness between metrics scrapes. least-load routes to the




Figure 2: Proxy-to-replica RTT over the Apr 5 16:15–16:45 tuning window. Bold lines show the EWMA-smoothed signal ($\alpha=0.3$) that GORGO uses for routing; faint lines show raw probe samples. Dashed lines show per-replica means (Seoul 456 ms, Frankfurt 279 ms, Ashburn 37 ms). Seoul swings widely from trans-Pacific routing-table churn; Frankfurt is moderately variable; Ashburn is near-constant. The wide cross-region spread motivates network-aware routing.

149 replica with the lowest aggregate token-weighted load, summing running and queued requests with
 150 used and queued KV tokens; it is heavier-signal than least-request but cache-blind. `prefix-cache`
 151 is the longest-prefix-match policy of Preble [Srivatsa et al., 2025] as packaged in AIBrix [AIB-
 152 rix Team, 2024]: it picks the replica with the longest cached prefix match and falls back to least-
 153 request when the running-request imbalance exceeds a soft threshold. `simple-session-affinity`
 154 hashes the first 256 token IDs of the prompt modulo the number of replicas, maximizing intra-
 155 user reuse but providing no load-escape mechanism. We evaluate these six policies in each
 156 of two windows; all six fleets start simultaneously at the same wall-clock time with the same
 157 input trace. W1 supplies `gorgo-hillclimb` with uniform starting weights (`rtt_weight=1.0`,
 158 `prefill_weight=1.0`, `load_weight=1.0`); W2 uses the weights that `gorgo-hillclimb` learned dur-
 159 ing W1 (`rtt_weight`≈0.39, `prefill_weight`≈1.88, `load_weight`≈6.38). `gorgo-static` in W2 de-
 160 ploys those weights without further adjustment. Concurrency is 64 in-flight requests per proxy.

161 **Metrics.** For each policy and window we report TTFT p50, p95, and p99 over all successful
 162 streaming responses. Traces are pre-filtered to the model’s effective context boundary; all policies
 163 achieve 100% success rate in all windows.

164 5 Results

165 5.1 p95 TTFT Objective

166 **Results across three windows.** Table 2 reports all three decoded-trace windows under the frozen
 167 2D weights ($w_{\text{rtt}}=0.276$, $w_{\text{queue}}=0.5$). On the in-sample Apr 5 tuning window—at full arrival rate,
 168 where every policy is saturated—GORGO wins TTFT p50 and p99 and cuts E2E p95 by 34.9% over
 169 the best baseline (8.91 s vs. 13.69 s), trailing `simple-session-affinity` by only 3.5% on TTFT
 170 p95. On both held-out windows—Apr 6 at half load and Apr 7 at third load, each within capacity
 171 for all five policies—GORGO *sweeps every metric*: on Apr 6, TTFT p95 is 15.5% better than
 172 `simple-session-affinity` (1,584 vs. 1,875 ms) and E2E p95 is 30.8% better (3.29 vs. 4.75 s);
 173 on Apr 7, TTFT p95 is 7.9% better (1,377 vs. 1,495 ms) and E2E p95 is 14.5% better (3.34 vs.
 174 3.90 s). The weights, learned once on Apr 5, transfer to later days and times of day without re-
 175 tuning. Figure 3 shows the tuning run that produced the weights.

176 5.2 p50 TTFT Objective

177 *[Reviewer comment: This p50 subsection and the Objective-Sensitivity table (Table 6) still report*
 178 *the OLD 3-weight model on the apr2 windows ($w_{\text{rtt}}/w_{\text{prefill}}/w_{\text{load}} = 0.39/1.88/6.38$), which is*
 179 *inconsistent with the 2D results above (2 free weights $w_{\text{rtt}}, w_{\text{queue}}$ on decoded Apr 5–7). To recon-*

Table 2: Decoded-trace results across three 30-minute windows. GORGO uses the 2D cost model with weights $w_{\text{rtt}}=0.276$, $w_{\text{queue}}=0.5$, learned by (1+1)-ES on the Apr 5 tuning window and *frozen* for both held-out windows. TTFT in ms, E2E p95 in s; bold is best in column within each window. The held-out windows are replayed within capacity for every policy (scheduling-slip p95 ≤ 10 ms).

Policy	TTFT p50 (ms)	TTFT p95 (ms)	TTFT p99 (ms)	E2E p95 (s)
<i>Tuning window — Apr 5 16:15–16:45 (in-sample, full load; $n=7,195$)</i>				
GORGO	673	2,514	4,473	8.91
simple-session-affinity	712	2,428	5,190	18.27
least-load	763	2,447	4,349	13.69
least-request	968	3,970	7,012	16.62
prefix-cache	1,477	6,784	9,613	22.63
<i>Held-out eval — Apr 6 15:05–15:35, half load ($n=7,630$)</i>				
GORGO	491	1,584	2,222	3.29
simple-session-affinity	627	1,875	3,056	4.75
least-load	616	1,818	2,885	4.06
least-request	634	1,852	2,484	3.99
prefix-cache	574	1,798	2,750	4.95
<i>Held-out eval — Apr 7 19:45–20:15, third load ($n=8,663$)</i>				
GORGO	386	1,377	1,973	3.34
simple-session-affinity	439	1,495	2,313	3.90
least-load	480	1,637	2,294	4.04
least-request	509	1,724	2,485	4.40
prefix-cache	516	1,830	2,801	11.98




Figure 3: (1+1)-ES convergence on the Apr 5 tuning window (p95 TTFT objective). *Top-left*: best-incumbent weight trajectories— w_{queue} rails at its 0.5 ceiling (maximal load avoidance) while w_{rtt} settles near 0.28; dots show ES proposals. *Top-right*: step size σ rises then decays as the search converges. *Bottom-left*: objective (negative p95 TTFT) climbs from -5.9 to -1.2 . *Bottom-right*: acceptance rate falls toward the 1/5 target.

180 *cile: re-run the p50 objective under the 2D model, or relocate this to an appendix as the 3-weight*
181 *precursor. The real-trace numbers here are Modal-only (not locally verifiable).]*

182 To test whether the learned weights depend on the objective percentile, we re-run W1 tuning with a
183 p50 TTFT objective (same trace, same hyperparameter ranges).

184 The p50-tuned ES converged to $w_{\text{rtt}}=1.21$, $w_{\text{prefill}}=0.52$, $w_{\text{load}}=1.69$ —a qualitatively different
185 operating point from the p95 run. gorgo-hillclimb achieves 0.163 s p50 TTFT (16.8% gap over
186 simple-session-affinity) and 0.943 s p95 TTFT (4.5% gap). However, the RTT-first routing
187 concentrates traffic on the nearest replica, degrading E2E: p95 E2E is 2.40 s, 14.1% worse than

Table 3: ARTChat-411k W1 with p50 TTFT objective (tuning window, 00:30–01:00 UTC, April 2nd 2026). Same trace and ranges as the p95 run. $n=2,095$ requests per policy; 100% success rate.

Policy	TTFT (s)			E2E (s)			Succ. %
	p50	p95	p99	p50	p95	p99	
gorgo-hillclimb	0.163	0.943	2.185	1.45	2.40	3.95	100
simple-session-affinity	0.196	0.988	1.654	1.48	2.61	3.47	100
least-request	0.207	1.102	1.934	1.24	2.10	3.07	100
least-load	0.264	1.079	1.874	1.40	2.38	2.92	100
random	0.274	1.186	2.021	1.27	2.33	3.24	100
prefix-cache	0.287	1.052	1.773	1.31	2.34	3.32	100

Table 4: p50-objective W2a (held-out nighttime, 01:00–01:30 UTC, April 2nd 2026). gorgo-static deploys the p50-tuned weights ($w_{\text{rtt}}=1.21$, $w_{\text{prefill}}=0.52$, $w_{\text{load}}=1.69$). $n=2,207$ requests per policy; 100% success rate.

Policy	TTFT (s)			E2E (s)			Succ. %
	p50	p95	p99	p50	p95	p99	
gorgo-static	0.157	0.961	1.776	1.42	2.42	3.27	100
simple-session-affinity	0.206	0.995	1.960	1.56	2.67	3.45	100
least-request	0.238	1.180	1.812	1.25	2.27	3.27	100
least-load	0.289	1.298	2.052	1.41	2.57	3.27	100
prefix-cache	0.290	1.160	1.944	1.51	2.51	8.65	100

least-request (2.10 s). This tradeoff—better median TTFT at the cost of E2E—is the mirror image of the p95 policy, which wins both.

Held-out evaluation. In W2a, gorgo-static with p50-tuned weights sweeps TTFT: p50 (0.157 s, 23.8% gap over simple-session-affinity), p95 (0.961 s, 3.4% gap), and p99 (1.776 s). E2E p95 is 2.42 s—6.6% worse than least-request (2.27 s), consistent with the RTT-first routing tradeoff observed in W1.

In W2b (midday), gorgo-static with p50-tuned weights wins TTFT p50 (0.188 s, 21.7% gap), p95 (1.184 s, 12.0% gap over least-request), and p99 (1.901 s, 12.3% gap). E2E p95 is 2.92 s, 8.9% worse than least-request (2.68 s). The p50-tuned policy achieves stronger TTFT p50 margins than the p95-tuned policy (21.7% vs. 16.4%) but pays a consistent E2E penalty—the lower $w_{\text{load}}=1.69$ (vs. 6.38) provides less load balancing under the heavier midday traffic.

5.3 Objective Sensitivity: Weight Comparison

Table 6 compares the learned weights across objectives.

The p95 objective drives the policy toward cache-first routing ($w_{\text{prefill}}=1.88$) with aggressive load avoidance ($w_{\text{load}}=6.38$), because tail requests are the ones stuck behind queues on cache-cold replicas. The p50 objective drives toward RTT-first routing ($w_{\text{rtt}}=1.21$, $3.1\times$ larger), because the typical request is short enough that network latency dominates over cache effects. The p50-tuned policy achieves 0.163 s p50 TTFT (vs. 0.180 s for p95-tuned), a 9% improvement on the metric it optimizes.

6 Limitations

p99 noise floor. Each window has $\approx 2,100$ – $2,200$ successful responses, so the p99 standard error is $\sigma_{\text{TTFT}}/\sqrt{n \cdot 0.01 \cdot 0.99} \approx 0.07$ s. *[Reviewer comment: Stale: the 1.626/1.702 values below match no table and refer to the apr2 W2 window, which the p95 results no longer use (now the decoded Apr 5–7 windows); reconcile or remove once p50/p99 are migrated to the 2D results.]* The W2 p99 gap between gorgo-static (1.626 s) and the next-best policy prefix-cache (1.702 s) is 0.076 s—just above one SE—so the p99 win, while directionally consistent with p50 and p95, is not individually significant at this sample size.

Table 5: p50-objective W2b (held-out midday diurnal, 12:30–13:00 UTC, April 2nd 2026). $n=3,071$ requests per policy; 100% success rate.

Policy	TTFT (s)			E2E (s)			Succ. %
	p50	p95	p99	p50	p95	p99	
gorgo-static	0.188	1.184	1.901	1.60	2.92	4.83	100
simple-session-affinity	0.240	1.498	2.523	1.71	3.78	8.21	100
random	0.289	1.421	2.331	1.40	2.93	3.97	100
prefix-cache	0.293	1.451	2.973	1.59	3.22	5.29	100
least-request	0.293	1.345	2.167	1.36	2.68	3.83	100
least-load	0.303	1.831	3.144	1.66	3.91	8.36	100

Table 6: Learned weights under different objective metrics, same tuning trace and hyperparameter ranges. The ES discovers qualitatively different operating points depending on which percentile it optimizes.

Objective	w_{rtt}	w_{prefill}	w_{load}
neg_p95_ttft	0.39	1.88	6.38
neg_p50_ttft	1.21	0.52	1.69

215 **Scope of fleet.** Each policy runs on its own dedicated three-replica fleet of identical GPUs. Multi-
 216 tenant sharing, heterogeneous hardware, and prefill/decode disaggregation [Zhong et al., 2024, Patel
 217 et al., 2024] are not characterized here and would require additional cost terms. The score uses
 218 only HTTP-exposed metrics, not scheduler-internal signals [Pan et al., 2025, Yang et al., 2025a] that
 219 could tighten the prefill estimate.

220 **Live-tuning instability.** Running the (1+1)-ES live can produce a heavy-tailed p99 excursion
 221 because the negative-p95 objective does not penalize a single bad-minute proposal fast enough. The
 222 freeze-for-production recommendation in §5.1 is therefore a safety property, not just a convenience.
 223 Decode dynamics such as the effect of variable length decode and high memory consumption on
 224 load may impact live-tuning instability; exploring the impact of higher output token limits, or more
 225 sophisticated online or machine learning algorithms, would be a direction for future work.

226 7 Related Work

227 **Prefix caches and reuse.** RadixAttention in SGLang [Zheng et al., 2024b] stores per-request KV
 228 state in a radix tree; PagedAttention in vLLM [Kwon et al., 2023] enables prefix sharing through
 229 paged memory. KVLink [Yang et al., 2025b], ChunkKV [Liu et al., 2025], KVFlow [Pan et al.,
 230 2025], and Learned Prefix Caching [Yang et al., 2025a] improve the single-replica cache but do not
 231 decide which replica serves a request.

232 **Cross-replica routing and cross-region serving.** Preble [Srivatsa et al., 2025] introduces longest-
 233 prefix-match routing across replicas with a load-balance fallback; AIBrix [AIBrix Team, 2024] pack-
 234 ages a production-grade variant of the same design and supplies the baseline we run. Mooncake [Qin
 235 et al., 2024] is a KV-cache-centric architecture and the source of our trace format. SkyServe [Mao
 236 et al., 2025] and SkyWalker [Xia et al., 2025] address cross-region LLM serving at the placement
 237 and spillover layers respectively, but treat per-request routing as out of scope. k-LPM [Dexter et al.,
 238 2025] formulates LLM scheduling under TTFT constraints as NP-hard. DLPM [Cao et al., 2025]
 239 targets fairness with locality. Lumnix [Sun et al., 2024] migrates requests across replicas. Multi-
 240 LLM routers [Wu and Silwal, 2025] address model selection. CacheBlend [Yao et al., 2024] extends
 241 KV reuse beyond exact prefixes by selectively recomputing cache for fused non-prefix chunks, but
 242 targets single-node reuse rather than cross-replica routing. DistServe [Zhong et al., 2024] and Split-
 243 wise [Patel et al., 2024] disaggregate prefill from decode. GORGON is the first single-router policy in
 244 this space that scores cache locality, replica load, and wide-area latency in one cost and derives the
 245 cost weights from the deployment’s own per-request TTFT stream, with no engine modifications.

246 **Online hyperparameter adaptation.** Bayesian and bandit-based methods have been applied to
 247 serving configuration [Wu and Silwal, 2025]. For a 2-dimensional parameter space the (1+1)-ES
 248 [Rechenberg, 1973] provides fast, overhead-free convergence with no surrogate model.

249 8 Conclusion

250 Routing across LLM replicas is a three-signal decision balancing cache locality, replica load, and
 251 wide-area latency, but production heuristics commit to one signal and degrade when that stops being
 252 the limiting factor. We treat the three as terms of an additive per-replica cost and let online TTFT
 253 feedback optimize scaling weights, in place of operator-tuned constants or offline profiling. On
 254 a long-context, high-prefix-reuse production trace the GORGO policy family collectively achieves
 255 the lowest TTFT at every reported percentile in both a tuning and a held-out evaluation windows.
 256 On a short-prompt, low-reuse public trace (Appendix A) the same policy is non-competitive, in
 257 agreement with the regime characterization in §3. The advantage is therefore a property of the
 258 workload regime as much as of the policy: it scales with how much of TTFT is recoverable through
 259 prefill reuse, queue avoidance, or network selection rather than fixed by hardware. The design
 260 choices of GORGO transfer to deployments where the workload regime is not known in advance
 261 and a dedicated calibration window before each redeploy is not affordable, making it effective in
 262 production LLM serving on long context, distributed workloads.

263 References

- 264 AIBrix Team. AIBrix: An open-source infrastructure for LLM inference at scale. <https://github.com/vllm-project/aibrix>, 2024.
 265
- 266 Shiyi Cao, Yichuan Wang, Ziming Mao, Pin-Lun Hsu, Liangsheng Yin, Tian Xia, Dacheng Li, Shu
 267 Liu, Yineng Zhang, Yang Zhou, Ying Sheng, Joseph Gonzalez, and Ion Stoica. Locality-aware
 268 fair scheduling in LLM serving. *arXiv preprint arXiv:2501.14312*, 2025.
- 269 Gregory Dexter, Shao Tang, Ata Fatahi, Qingquan Song, Tejas Dharamsi, and Aman Gupta. LLM
 270 query scheduling with prefix reuse and latency constraints. In *Advances in Neural Information
 271 Processing Systems (NeurIPS)*, 2025.
- 272 GLM Team. GLM-4: A family of open multilingual multimodal chat LLMs, 2024.
- 273 Yiyuan He, Minxian Xu, Jingfeng Wu, Jianmin Hu, Chong Ma, Min Shen, Le Chen, Chengzhong
 274 Xu, Lin Qu, and Kejiang Ye. BanaServe: Unified KV cache and dynamic module migration for
 275 balancing disaggregated LLM serving in AI infrastructure, 2025.
- 276 David Karger, Eric Lehman, Tom Leighton, Rina Panigrahy, Matthew Levine, and Daniel Lewin.
 277 Consistent hashing and random trees: Distributed caching protocols for relieving hot spots on the
 278 World Wide Web. In *Proceedings of the Twenty-Ninth Annual ACM Symposium on Theory of
 279 Computing (STOC)*, pages 654–663, 1997. doi: 10.1145/258533.258660.
- 280 Woosuk Kwon, Zhuohan Li, Siyuan Zhuang, Ying Sheng, Lianmin Zheng, Cody Hao Yu, Joseph E.
 281 Gonzalez, Hao Zhang, and Ion Stoica. Efficient memory management for large language model
 282 serving with PagedAttention. In *ACM Symposium on Operating Systems Principles (SOSP)*, 2023.
- 283 Xiang Liu, Zhenheng Tang, Peijie Dong, Zeyu Li, Yue Liu, Bo Li, Xuming Hu, and Xiaowen Chu.
 284 ChunkKV: Semantic-preserving KV cache compression for efficient long-context LLM inference.
 285 In *Advances in Neural Information Processing Systems (NeurIPS)*, 2025.
- 286 Ziming Mao, Tian Xia, Zhanghao Wu, Wei-Lin Chiang, Tyler Griggs, Romil Bhardwaj, Zongheng
 287 Yang, Scott Shenker, and Ion Stoica. SkyServe: Serving ai models across regions and clouds
 288 with spot instances. In *Proceedings of the Twentieth European Conference on Computer Systems
 289 (EuroSys ’25)*, 2025. doi: 10.1145/3689031.3717459.
- 290 Zaifeng Pan, Ajikumar Patel, Zhengding Hu, Yipeng Shen, Yue Guan, Wan-Lu Li, Lianhui Qin,
 291 Yida Wang, and Yufei Ding. KVFlow: Efficient prefix caching for accelerating LLM-based
 292 multi-agent workflows. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2025.
 293 arXiv:2507.07400.

294 Pratyush Patel, Esha Choukse, Chaojie Zhang, Aashaka Shah, Íñigo Goiri, Saeed Maleki, and
 295 Ricardo Bianchini. Splitwise: Efficient generative LLM inference using phase splitting. In
 296 *ACM/IEEE International Symposium on Computer Architecture (ISCA)*, 2024.

297 Ruoyu Qin, Zheming Li, Weiran He, Mingxing Zhang, Yanzhe Wu, Weimin Zheng, Xinran Chen,
 298 Rui Xu, Xingda Yao, Rongxin Wei, Zhiyuan Zhao, Jing Huang, Mingjie Yang, and Jidong Wu.
 299 Mooncake: A KVCache-centric disaggregated architecture for LLM serving, 2024.

300 Qwen Team. Qwen3 technical report. 2025. Alibaba Cloud.

301 Ingo Rechenberg. Evolutionsstrategie: Optimierung technischer systeme nach prinzipien der biolo-
 302 gischen evolution. 1973.

303 Vikranth Srivatsa, Sumanth Madhavan, Tian Hu, Sarah Tarun, Yuhan Cheng, Wei Han, Jingyu Yu,
 304 Roberto Cristian Rusti, Lifeng Wang, Yiyang Liu, and Hao Zhang. Preble: Efficient distributed
 305 prompt scheduling for LLM serving. In *International Conference on Learning Representations*
 306 *(ICLR)*, 2025.

307 Ion Stoica, Robert Morris, David Liben-Nowell, David R. Karger, M. Frans Kaashoek, Frank Dabek,
 308 and Hari Balakrishnan. Chord: A scalable peer-to-peer lookup protocol for internet applications.
 309 *IEEE/ACM Transactions on Networking*, 11(1):17–32, 2003. doi: 10.1109/TNET.2002.808407.

310 Biao Sun, Ziming Huang, Hanyu Tang, Chengquan Wang, Cheng Zhang, Yiming Li, Liu Yang,
 311 Wencong Zhang, Lin-Wei Wang, Tianhe Wu, Ang Cao, Yongqiang Lin, et al. Llumnix: Dynamic
 312 scheduling for large language model serving. In *USENIX Symposium on Operating Systems De-
 313 sign and Implementation (OSDI)*, 2024.

314 Fangzhou Wu and Sandeep Silwal. Efficient training-free online routing for high-volume multi-LLM
 315 serving. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2025.

316 Tian Xia, Ziming Mao, Jamison Kerney, Ethan J. Jackson, Zhifei Li, Jiarong Xing, Scott Shenker,
 317 and Ion Stoica. SkyWalker: A locality-aware cross-region load balancer for LLM inference. arXiv
 318 preprint arXiv:2505.24095, 2025.

319 Dongsheng Yang, Austin Li, Kai Li, and Wyatt Lloyd. Learned prefix caching for efficient LLM
 320 inference. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2025a.

321 Jingbo Yang, Bairu Hou, Wei Wei, Yujia Bao, and Shiyu Chang. KVLink: Accelerating large
 322 language models via efficient KV cache reuse. *arXiv preprint arXiv:2502.16002*, 2025b.

323 Jiayi Yao, Hanchen Li, Yuhan Liu, Siddhant Ray, Yihua Peng, Hao Zhang, Ion Stoica, Kurt Keutzer,
 324 and Michael W. Mahoney. CacheBlend: Fast large language model serving for RAG with cached
 325 knowledge fusion, 2024.

326 Gyeong-In Yu, Jinseok Jeong, Gyeongchan Kim, Suhyun Kim, and Byung-Gon Chun. Orca: A
 327 distributed serving system for transformer-based generative models. In *Proceedings of the 16th*
 328 *USENIX Symposium on Operating Systems Design and Implementation (OSDI 22)*, 2022.

329 Wenting Zhao, Xiang Ren, Jack Hessel, Claire Cardie, Yejin Choi, and Yuntian Deng. WildChat: 1M
 330 ChatGPT interaction logs in the wild. In *International Conference on Learning Representations*
 331 *(ICLR)*, 2024.

332 Lianmin Zheng, Wei-Lin Chiang, Ying Sheng, Tianle Li, Siyuan Zhuang, Zhanghao Wu, Yong-
 333 hao Zhuang, Zhuohan Li, Zi Lin, Eric P. Xing, Joseph E. Gonzalez, Ion Stoica, and Hao Zhang.
 334 LMSYS-Chat-1M: A large-scale real-world LLM conversation dataset. In *International Confer-
 335 ence on Learning Representations (ICLR)*, 2024a.

336 Lianmin Zheng, Liangsheng Yin, Zhiqiang Xie, Chuyue Livia Sun, Jeff Huang, Cody Hao Yu,
 337 Shiyi Cao, Christos Kozyrakis, Ion Stoica, Joseph E. Gonzalez, Clark Barrett, and Ying Sheng.
 338 SGLang: Efficient execution of structured language model programs. In *Advances in Neural*
 339 *Information Processing Systems (NeurIPS)*, 2024b.

340 Yinmin Zhong, Shengyu Liu, Junda Chen, Jianbo Hu, Yibo Zhu, Xuanzhe Liu, Xin Jin, and Hao
 341 Zhang. DistServe: Disaggregating prefill and decoding for goodput-optimized large language
 342 model serving. In *USENIX Symposium on Operating Systems Design and Implementation (OSDI)*,
 343 2024.

NeurIPS Paper Checklist

1. Claims

Question: Do the main claims made in the abstract and introduction accurately reflect the paper’s contributions and scope?

Answer: [Yes]

Justification: The abstract claims (i) public datasets underrepresent long-context multi-turn traffic (quantified in §3), (ii) we evaluate routing policies on a production trace (Table 2), and (iii) GORGO sweeps every TTFT percentile on the held-out windows. We are explicit in the introduction, §5.1, and §6 that “GORGO” refers to a family, that the W2 p99 margin is at the noise floor, and that on the daytime stress window gorgo-hillclimb’s p99 inflates and prefix-cache wins p99 instead.

2. Limitations

Question: Does the paper discuss the limitations of the work?

Answer: [Yes]

Justification: Section 6 lists single-trace generalization, the W2 p99 noise floor, the single-tenant assumption, gorgo-autotune instability, hardware homogeneity, the absence of PD-disaggregation, and ES convergence time.

3. Theory assumptions and proofs

Question: For each theoretical result, does the paper provide the full set of assumptions and a complete proof?

Answer: [N/A]

Justification: The paper presents an additive cost model and an online optimizer with no formal theorems.

4. Experimental result reproducibility

Question: Does the paper fully disclose all the information needed to reproduce the main experimental results?

Answer: [Yes]

Justification: Section 4 specifies model, regions, hardware, fleet topology, concurrency, trace windows, and per-window GORGO seeding. Section 2 gives Equation 3, the (1+1)-ES configuration ($\sigma_0 = 0.5$, 1/5-rule, hop 16, window 64), and the median-of-rates fit.

5. Open access to data and code

Question: Does the paper provide open access to the data and code?

Answer: [Yes]

Justification: The proxy and policy code are available at <https://anonymous.4open.science/r/GORGO-D25A>. The ARTChat-411k trace is not redistributable under its terms of service. The public-dataset characterization (LMSYS, WildChat) is fully reproducible with no proprietary inputs.

6. Experimental setting/details

Question: Does the paper specify all the training and test details?

Answer: [Yes]

Justification: Section 4 specifies model, hardware, regions, concurrency, trace windows, max input tokens, and per-window seeding for adaptive policies. ES configuration is given in Section 2.

7. Experiment statistical significance

Question: Does the paper report error bars or appropriate statistical significance information?

Answer: [Yes]

Justification: Section 6 derives an approximate p99 standard error and shows that the W2 p99 margin between gorgo-hillclimb and random is inside one standard error, while the W1 sweep and W2 p50/p95 wins are several standard errors clear of zero.

8. Experiments compute resources

Question: Does the paper provide sufficient information on compute resources?

Answer: [Yes]

Justification: Each policy runs on a dedicated 3-replica fleet of L40S SGLang servers (tensor-parallel 2, 96 GB VRAM per replica). Eight policies $\times$ two windows = 16 fleet-runs, each ~ 30 min wall clock, on three regions. Total GPU-hours: ≈ 24 .

400 **9. Code of ethics**
 401 Question: Does the research conform with the NeurIPS Code of Ethics?
 402 Answer: [Yes]
 403 Justification: The work concerns routing infrastructure. The production trace is used un-
 404 der terms that permit research replay; we release only aggregated prefix-trie statistics, not
 405 individual requests. No individuals are identifiable from the released aggregates.

406 **10. Broader impacts**
 407 Question: Does the paper discuss potential societal impacts?
 408 Answer: [Yes]
 409 Justification: Routing improvements reduce the energy and dollar cost of serving LLMs at
 410 fixed quality. Lower-cost serving may accelerate deployment in settings where that is so-
 411 cially harmful; the contributions are infrastructure-level and do not affect model capability.

412 **11. Safeguards**
 413 Question: Does the paper describe safeguards for responsible release?
 414 Answer: [N/A]
 415 Justification: We do not release datasets or models. The released code is a routing proxy
 416 and policy library with no inherent dual-use beyond LLM-serving infrastructure.

417 **12. Licenses**
 418 Question: Are creators or owners of assets properly credited?
 419 Answer: [Yes]
 420 Justification: SGLang, vLLM, AIBrix, LMSYS-Chat-1M, WildChat-4.8M, and Qwen3 are
 421 cited and used under their published licenses.

422 **13. New assets**
 423 Question: Are new assets introduced in the paper well documented?
 424 Answer: [Yes]
 425 Justification: The proxy, policy library, experiment controller, and trace builder are doc-
 426 umented in the repository’s top-level documentation. Spec and manifest files are stored
 427 alongside the controller.

428 **14. Crowdsourcing and research with human subjects**
 429 Question: Does the paper involve crowdsourcing or human subjects?
 430 Answer: [N/A]
 431 Justification: No human subjects.

432 **15. IRB approvals**
 433 Answer: [N/A]
 434 Justification: No human subjects.

435 **16. Declaration of LLM usage**
 436 Question: Does the paper describe the usage of LLMs if it is an important, original, or
 437 non-standard component of the research methodology?
 438 Answer: [N/A]
 439 Justification: LLM serving infrastructure is the object of study. Authors used LLMs for
 440 routine writing assistance only.

A WildChat replay (regime-sensitivity control)

We replay a 30-minute window of WildChat-4.8M [Zhao et al., 2024] on the same $2 \times L40S$ three-region fleet, with the same eight policies as the ARTChat-411k sweeps. WildChat averages 2,925 input tokens per request and 5.3% intra-user prefix reuse—an order of magnitude shorter than ARTChat-411k and roughly a tenth the reuse—putting it well outside the regime characterized in §3.

Table 7: WildChat-4.8M replay: 30-minute window, $c=32$. TTFT in seconds. Bold is best in column. Compared to the ARTChat-411k results, GORGO variants are not competitive: gorgo-static and gorgo-autotune are the worst two policies on p95.

Policy	Completed	p50	p95	p99	max	Succ. %
simple-session-affinity	20,447	0.416	1.025	1.712	7.75	99.6
least-request	20,517	0.480	1.036	1.731	110.45	100.0
random	20,436	0.416	1.077	1.796	15.82	99.6
prefix-cache	20,504	0.497	1.186	2.588	10.63	99.9
least-load	20,484	0.532	1.217	2.535	8.28	99.8
gorgo-hillclimb	20,473	0.541	1.325	2.654	7.05	99.8
gorgo-static	20,462	0.709	1.932	3.969	60.18	99.7
gorgo-autotune	20,451	0.541	3.583	6.438	12.55	99.7




Figure 4: WildChat-4.8M replay, per-percentile TTFT (p50/p95/p99) across the eight policies. GORGO variants (dark blues) are not competitive: gorgo-static and gorgo-autotune are the slowest two policies at the tail, and gorgo-hillclimb sits mid-pack. The ranking inverts the ARTChat-411k result, consistent with the regime argument in §3.

The ranking is reversed: simple-session-affinity wins p95 and p99 by a large margin, least-request is competitive, gorgo-hillclimb sits mid-pack, and gorgo-static and gorgo-autotune are the worst two policies. Two mechanisms drive this. First, the prefill term in Equation 3 is small at 2,925 tokens; the cost-model weights are mostly fitting noise. Second, intra-user reuse is 5.3% rather than 53%; even a perfect cache-routing decision saves only a few hundred tokens per request, well inside the cross-region RTT band, so chasing cache locality is counterproductive when paired with a cross-region routing penalty. The gorgo-static weights frozen from an ARTChat-411k W1 are particularly mis-calibrated to this regime—a reminder that the cost-model parameters do not transfer across workload classes.

We include this result as a regime-sensitivity control, not as a contribution. The takeaway is the same one stated in §3: the gain from joint cache-network-queue routing depends on the workload having long prompts and substantial intra-user prefix reuse. Public chat datasets do not exercise this regime.




Figure 5: Achieved prefix hit rate over the W1 tuning window: the fraction of each request’s input tokens that were already cached on the replica the policy routed to (not the total cache available across all replicas). A higher rate means the routing decision exploited more of the available cache. Bold lines are EWMA-smoothed ($\alpha=0.06$); faint lines are 64-request rolling averages. gorgo-hillclimb converges from $\sim 45\%$ to $\sim 95\%$ within the first 5 minutes as the ES learns to route requests to their best-cached replica.

460 B GORGO Achieves Convergence of Cache Reuse

461 C Load-Weight Ablation: Single-Replica Concentration

462 When w_{load} is unconstrained (range $[10^{-5}, 5]$), the ES drives it to zero: the resulting policy routes
 463 100% of traffic to a single replica (Figure 6). On the nighttime tuning trace (~ 1.2 req/s) this pro-
 464 duces the best TTFT because every request hits a warm cache on the closest replica with no load
 465 penalty. On the heavier midday diurnal trace (~ 1.7 req/s, 47% more requests), the single replica
 466 saturates: decode throughput drops to 70 tok/s (vs. 119 tok/s with $w_{\text{load}}=6.38$) and E2E p95 inflates
 467 to 12.58 s— $4.8\times$ worse than least-request.

468 Constraining the ES search range to $w_{\text{load}} \in [0.1, 10.0]$ prevents this degenerate solution. The ES
 469 converges to $w_{\text{load}} \approx 6.38$, which actively avoids loaded replicas and achieves the best TTFT *and*
 470 E2E on both evaluation windows (§5.1). Table 8 compares the two configurations on the midday
 471 diurnal trace.

Table 8: Midday diurnal evaluation: unconstrained ($w_{\text{load}}=0$) vs. constrained ($w_{\text{load}}=6.38$) gorgo-static. Constraining w_{load} trades 3% TTFT p95 for $5\times$ better E2E p95.

	TTFT p50	TTFT p95	E2E p50	E2E p95	Decode	Concentration
$w_{\text{load}}=0$	0.190	1.101	2.07	12.58	70	100%
$w_{\text{load}}=6.38$	0.195	1.136	1.29	2.46	119	60%
Δ	+2.6%	+3.2%	−37.7%	−80.4%	+70%	




Figure 6: *Left*: TTFT p95 (dark) vs. E2E p95 (light) on the midday diurnal trace with $w_{load}=0$. *gorgo-static* wins TTFT but its E2E inflates to 12.6 s. *Right*: routing concentration. *gorgo-static* sends 100% of requests to one replica.